


Qwen2.5-1.5B: measured throughput by phase



tokens/s; seven llama.cpp repetitions. Prefill uses 2048 tokens; decode uses 128.